

Spatiotemporal Graph Deep Learning



Traffic monitoring



Smart cities



Energy analytics



Physics




Stock markets

Outline

- 1) Spatiotemporal time series
- 2) Spatiotemporal GNNs
- 3) Dealing with missing data

 demo

Deep learning for time series forecasting




The standard deep learning approach to time series forecasting consists in training a **single neural network** on a collection of **time series**.

- Each time series is treated **independently** from the others.
- A single set of **shared** learnable **parameters** is used to predict each time series.
- Resulting models are **effective** and **efficient**.

! **Dependencies** across time series are often **discarded**.

[1] K. Benidis *et al.*, “Deep Learning for Time Series Forecasting: Tutorial and Literature Survey”, ACM CS 2022.

Relational inductive biases



One way out is to embed such **relational structure** as an **architectural bias** into the processing.

Graph neural networks provide appropriate neural operators.

- **Message-passing** blocks allow for **localizing** the **predictions**
→ conditioning on observations at related time series (neighboring nodes).
- **Parameters** are **shared** and the model can operate on arbitrary sets of time series.

[2] D. Bacciu *et al.*, “A gentle introduction to deep learning for graphs”, NN 2020.


[3] M. M. Bronstein *et al.*, “Geometric deep learning: Grids, groups, graphs, geodesics, and gauges” 2021.

Spatiotemporal time series

Collections of time series

We consider a set of N **correlated time series**, where each i -th time series is associated with:

- an **observation vector** $\mathbf{x}_t^i \in \mathbb{R}^{d_x}$ at each time step t ;
- a vector of **exogenous variable** $\mathbf{u}_t^i \in \mathbb{R}^{d_u}$ at each time step t ;
- a vector of **static (time-independent) attributes** $\mathbf{v}^i \in \mathbb{R}^{d_v}$.



Capital letters denote the stacked representations encompassing the N time series in the collection, e.g., $\mathbf{X}_t \in \mathbb{R}^{N \times d_x}$, $\mathbf{U}_t \in \mathbb{R}^{N \times d_u}$.

Correlated time series

We assume a **time-invariant** stochastic process


$$\mathbf{x}_t^i \sim p^i(\mathbf{x}_t^i | \mathbf{X}_{<t}, \mathbf{U}_{\leq t}, \mathbf{V})$$

generating the data \mathbf{x}_t^i for all $i = 1 \dots N$ and $t \in \mathbb{N}$.

Note that the time series:

- can be generated by **different processes**,
- can **depend on others**,
- are assumed homogenous, synchronous, regularly sampled.

→ These assumptions can be **relaxed**



Notation:

$$\mathbf{X}_{t:t+T} = [\mathbf{X}_t, \dots, \mathbf{X}_{t+T-1}]$$

$$\mathbf{X}_{<t} = [\mathbf{X}_0, \dots, \mathbf{X}_{t-2}, \mathbf{X}_{t-1}]$$

Relational information



We assume the existence of **functional dependencies** between the time series.

- e.g., forecasts for one time series can be improved by accounting for the past values of other time series.

We model pairwise relationships existing at time step t with **adjacency matrix** $\mathbf{A}_t \in \{0, 1\}^{N \times N}$.

- \mathbf{A}_t can be **asymmetric** and **dynamic** (can vary with t).

- We call **spatial** the dimension spanning the time series collection.




Relational information with attributes

Optional **edge attributes** $e_t^{ij} \in \mathbb{R}^{d_e}$ can be associated to each non-zero entry of A_t .


The **set of attributed edges** encoding all the available relational information is denoted by

$$\mathcal{E}_t \doteq \{ \langle (i, j), e_t^{ij} \rangle \mid \forall i, j : A_t[i, j] \neq 0 \}.$$



→ For many applications, A_t **changes slowly** over time and can be considered as **constant** within a short window of observations.

Spatiotemporal time series



We use the terms **node** and **sensor** to indicate the N entities generating the time series.

→ We refer to the node set together with the relational information as **sensor network**.

The tuple $\mathcal{G}_t \doteq \langle \mathbf{X}_t, \mathbf{U}_t, \mathcal{E}_t, \mathbf{V} \rangle$ contain all the available information associated with time step t .

Example: Traffic monitoring system

Consider a sensor network monitoring the speed of vehicles at crossroads.



- $X_{<t}$ collects past **traffic** speed **measurements**.
- U_t stores identifiers for **time-of-the-day** and **day-of-the-week**.
- V collects static sensor's features, e.g., **type** or **number** of **lanes** of the monitored road.
- \mathcal{E} can be obtained by considering the **road network**.
 - Road closures and traffic diversions can be accounted for with a dynamic topology \mathcal{E}_t .

Time series forecasting


Multi-step time-series forecasting

Given a window of $W \geq 1$ **past** observations

$$\mathbf{X}_{t-W:t} = [\mathbf{X}_{t-W}, \dots, \mathbf{X}_{t-1}],$$

predict $H \geq 1$ **future** observations

$$\mathbf{x}_{t+h}^i, \quad i = 1 \dots N, h = 1 \dots H.$$



In particular, we are interested in learning a **parametric model** p_θ approximating the unknown data distribution p


$$p_\theta \left(\mathbf{x}_{t+h}^i \mid \mathbf{X}_{t-W:t}, \mathbf{U}_{t-W:t+h}, \mathbf{V} \right) \approx p^i \left(\mathbf{x}_{t+h}^i \mid \mathbf{X}_{<t}, \mathbf{U}_{\leq t+h}, \mathbf{V} \right).$$

- θ is the model parameter vector.

Time series forecasting + relational inductive biases

Condition the model on the relational information $\mathcal{E}_{t-W:t}$

$$p_{\theta}(\mathbf{x}_{t+h}^i | \mathcal{G}_{t-W:t}, \mathbf{U}_{t-W:t+h}, \mathbf{V})$$



⚡ The conditioning on the sequence of attributed graphs acts as a **regularization** to localize predictions w.r.t. the **neighborhood of each node**.

Point forecasts

For simplicity, we focus here on **point forecasts**, rather than the modeling of full data distributions p , and consider predictive model

$$\hat{\mathbf{x}}_{t+h}^i = \mathcal{F}(\mathcal{G}_{t-W:t}, \mathbf{U}_{t:t+h}; \boldsymbol{\theta})$$

where $\hat{\mathbf{x}}_{t+h}^i$ estimates $\mathbb{E}_p[\mathbf{x}_{t+h}^i]$

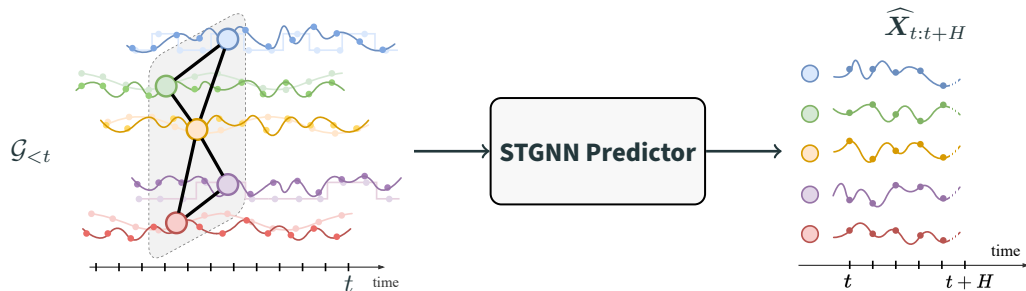
Parameters $\boldsymbol{\theta}$ can be learned by **minimizing a cost function** $\ell(\cdot, \cdot)$ (e.g., MSE) on a training set

$$\begin{aligned}\hat{\boldsymbol{\theta}} &= \arg \min_{\boldsymbol{\theta}} \frac{1}{NT} \sum_{t=1}^T \ell(\hat{\mathbf{X}}_{t:t+H}, \mathbf{X}_{t:t+H}) \\ &= \arg \min_{\boldsymbol{\theta}} \frac{1}{NT} \sum_{t=1}^T \left\| \mathbf{X}_{t:t+H} - \hat{\mathbf{X}}_{t:t+H} \right\|_2^2.\end{aligned}$$

Spatiotemporal Graph Neural Networks

Spatiotemporal Graph Neural Networks

We call **Spatiotemporal Graph Neural Network (STGNN)** a neural network exploiting both temporal and spatial relations of the input spatiotemporal time series.



We focus on models based on **message passing**.

Message-passing neural networks

To process the spatial dimension, we rely on the **message-passing (MP)** framework

$$\mathbf{h}^{i,l+1} = \text{UP}^l \left(\mathbf{h}^{i,l}, \text{AGGR}_{j \in \mathcal{N}(i)} \left\{ \text{MSG}^l(\mathbf{h}^{i,l}, \mathbf{h}^{j,l}, \mathbf{e}^{ji}) \right\} \right), \quad (1)$$


Where:

- $\text{MSG}^l(\cdot)$ is the **message function**, e.g., implemented by a MLP.
- $\text{AGGR}\{\cdot\}$ is the permutation invariant **aggregation function**.
- $\text{UP}^l(\cdot)$ is the **update function**, e.g., implemented by a MLP.

Aggregation is performed over $\mathcal{N}(i)$, i.e., the set of neighbors of node i .

[5] J. Gilmer *et al.*, “Neural message passing for quantum chemistry”, ICML 2017.

Message passing in action



Spatiotemporal message passing

Starting from the MP framework, we can define a general scheme for **spatiotemporal message-passing (STMP)** networks:

$$h_t^{i,l+1} = \text{UP}^l \left(h_{\leq t}^{i,l}, \text{AGGR}_{j \in \mathcal{N}_t(i)} \left\{ \text{MSG}^l(h_{\leq t}^{i,l}, h_{\leq t}^{j,l}, e_{\leq t}^{ji}) \right\} \right)$$

Rather than vectors, STMP blocks process **sequences**.

→ STMP blocks must be implemented with operators that work on sequences!

We will look at **different implementations** of STMP blocks in the following.

[4] A. Cini *et al.*, “Graph Deep Learning for Time Series Forecasting: A Comprehensive Methodological Framework” 2023.

A general recipe

STGNNs can be expressed as a sequence of three operations:

$$h_{t-1}^{i,0} = \text{ENCODER} \left(x_{t-1}^i, u_{t-1}^i, v^i \right), \quad (2)$$


$$H_{t-1}^{l+1} = \text{STMP}^l \left(H_{\leq t-1}^l, \mathcal{E}_{\leq t-1} \right), \quad l = 0, \dots, L-1 \quad (3)$$

$$\hat{x}_{t:t+H}^i = \text{DECODER} \left(h_{t-1}^{i,L}, u_{t:t+H}^i \right). \quad (4)$$

Where:

- $\text{ENCODER}(\cdot)$ is the encoding layer, e.g., implemented by a MLP.
- STMP is a stack of STMP layers.
- $\text{DECODER}(\cdot)$ is the readout layer, e.g., implemented by a MLP.

Framework overview



Design paradigms for STGNNs

Depending on the implementation of the STMP blocks, we categorize STGNNs into:

- **Time-and-Space (T&S)**

Temporal and spatial processing cannot be factorized in two separate steps.

- **Time-then-Space (TTS)**

Embed each time series in a vector, which is then propagated over the graph.

- **Space-then-Time (STT)**

Propagate nodes features at first and then process the resulting time series.



Time-and-Space

In T&S models, representations at every node and time step are the results of a **joint temporal and spatial encoding**

$$\mathbf{H}_{t-1}^{l+1} = \text{STMP}^l \left(\mathbf{H}_{\leq t-1}^l, \mathcal{E}_{\leq t-1} \right)$$

Several options exist.

- Integrate MP into neural operators for sequential data.
 - **Graph recurrent architectures, spatiotemporal convolutions, spatiotemporal attention, ...**
- Use temporal operators to compute messages.
 - **Temporal graph convolutions, spatiotemporal cross-attention, ...**
- Product graph representations.

Example 1: From Recurrent Neural Networks...

Consider a standard GRU [6] cell.

$$\mathbf{r}_t^i = \sigma \left(\Theta_r \left[\mathbf{x}_t^i || \mathbf{h}_{t-1}^i \right] + \mathbf{b}_r \right) \quad (5)$$

$$\mathbf{u}_t^i = \sigma \left(\Theta_u \left[\mathbf{x}_t^i || \mathbf{h}_{t-1}^i \right] + \mathbf{b}_u \right) \quad (6)$$

$$\mathbf{c}_t^i = \tanh \left(\Theta_c \left[\mathbf{x}_t^i || \mathbf{r}_t^i \odot \mathbf{h}_{t-1}^i \right] + \mathbf{b}_c \right) \quad (7)$$

$$\mathbf{h}_t^i = (1 - \mathbf{u}_t^i) \odot \mathbf{c}_t^i + \mathbf{u}_t^i \odot \mathbf{h}_{t-1}^i \quad (8)$$

Time series can be processed **independently** for each node or as a **single multivariate** time series.

[6] J. Chung *et al.*, “Empirical evaluation of gated recurrent neural networks on sequence modeling” 2014.

...to Graph Convolutional Recurrent Neural Networks

We can obtain a T&S model by implementing the gates of the GRU with MP blocks:

$$\mathbf{Z}_t^l = \mathbf{H}_t^{l-1} \quad (9)$$

$$\mathbf{R}_t^l = \sigma \left(\text{MP}_r^l \left([\mathbf{Z}_t^l || \mathbf{H}_{t-1}^l], \mathcal{E}_t \right) \right), \quad (10)$$

$$\mathbf{O}_t^l = \sigma \left(\text{MP}_o^l \left([\mathbf{Z}_t^l || \mathbf{H}_{t-1}^l], \mathcal{E}_t \right) \right), \quad (11)$$

$$\mathbf{C}_t^l = \tanh \left(\text{MP}_c^l \left([\mathbf{Z}_t^l || \mathbf{R}_t^l \odot \mathbf{H}_{t-1}^l], \mathcal{E}_t \right) \right), \quad (12)$$

$$\mathbf{H}_t^l = \mathbf{O}_t^l \odot \mathbf{H}_{t-1}^l + (1 - \mathbf{O}_t^l) \odot \mathbf{C}_t^l, \quad (13)$$

These T&S models are known as **graph convolutional recurrent neural networks (GCRNNs)** [7].

[7] Y. Seo *et al.*, “Structured sequence modeling with graph convolutional recurrent networks”, ICONIP 2018.

Popular GCRNNs

The first GCRNN has been introduced in [7], with MP blocks implemented as polynomial graph convolutional filters.

GCRNNs have become popular in the traffic forecasting context with the [Diffusion Convolutional Recurrent Neural Network \(DCRNN\)](#) architecture [8].

In DCRNN, MP is performed through [bidirectional diffusion convolution](#):

$$\mathbf{H}'_t = \sum_{k=0}^K \left(\mathbf{D}_{t,\text{out}}^{-1} \mathbf{A}_t \right)^k \mathbf{H}_t \mathbf{\Theta}_1^{(k)} + \left(\mathbf{D}_{t,\text{in}}^{-1} \mathbf{A}_t^\top \right)^k \mathbf{H}_t \mathbf{\Theta}_2^{(k)} \quad (14)$$

[7] Y. Seo *et al.*, “Structured sequence modeling with graph convolutional recurrent networks”, ICONIP 2018.

[8] Y. Li *et al.*, “Diffusion Convolutional Recurrent Neural Network: Data-Driven Traffic Forecasting”, ICLR 2018.

Example 2: Spatiotemporal convolutional networks (i)

A completely different approach is that of **spatiotemporal convolutional networks (STCNs)**, that **alternate spatial and temporal convolutional filters**:

- Compute intermediate representations by using a **node-wise temporal convolutional** layer:

$$z_{t-W:t}^{i,l} = \text{TCN}^l \left(h_{t-W:t}^{i,l-1} \right) \quad \forall i$$

where TCN^l indicates a temporal convolutional network layer.

- Then, compute the updated representation by using a **time-wise graph convolution**:

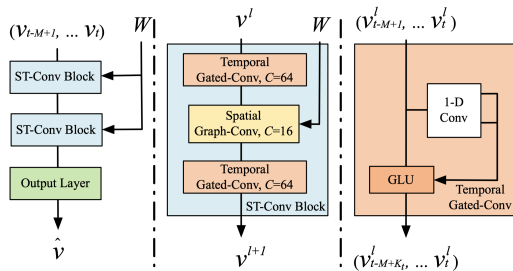
$$H_t^l = \text{MP}^l \left(Z_t^l, \mathcal{E}_t \right) \quad \forall t$$

Spatiotemporal convolutional networks (ii)

The first example of such architecture is the **STGCN** by Yu et al. [9].

The model is obtained by stacking STMP blocks consisting of

- a (gated) temporal convolution;
- a polynomial graph convolution;
- a second (gated) temporal convolution.



Courtesy of [9].

[9] B. Yu et al., “Spatio-temporal graph convolutional networks: a deep learning framework for traffic forecasting”, IJCAI 2018.

Example 3: Temporal Graph Convolution

A more integrated approach instead consists of **using temporal operators to compute messages**.

For example, we can design STMP layers s.t.

$$\mathbf{h}_{t-W:t}^{i,l} = \text{TCN}_1^l \left(\mathbf{h}_{t-W:t}^{i,l-1}, \text{AGGR}_{j \in \mathcal{N}_t(i)} \left\{ \text{TCN}_2^l \left(\mathbf{h}_{t-W:t}^{i,l-1}, \mathbf{h}_{t-W:t}^{j,l-1}, \mathbf{e}_{t-W:t}^{ji} \right) \right\} \right).$$

Analogous models can be built by exploiting attention-based operators [10], [11].


[10] I. Marisca *et al.*, “Learning to Reconstruct Missing Data from Spatiotemporal Graphs with Sparse Observations”, NeurIPS 2022.

[11] Z. Wu *et al.*, “TraverseNet: Unifying Space and Time in Message Passing for Traffic Forecasting”, TNNLS 2022.

Example 4: Product graph representations

Finally, an orthogonal option to those seen so far is to consider $\mathcal{G}_{t-W:t}$ as a **single spatiotemporal graph** \mathcal{S}_t .

Such **product graph** can be obtained by combining **temporal** and **spatial** graphs.




The resulting graph can be processed by any MP neural network.

[12] M. Sabbaqi *et al.*, “Graph-time convolutional neural networks: Architecture and theoretical analysis” 2022.

Building product graph representations


- **Cartesian product**

Spatial graphs are kept and each node is connected to itself in the previous time instant.



- **Kronecker product**

Each node is connected **only** to its neighbors in the previous time instant.

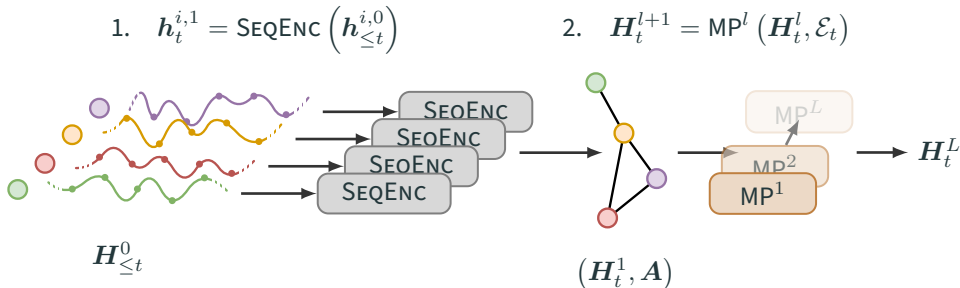


- ...


Time-then-Space models

The general recipe for a TTS model consists in:

1. **Embedding** each node-level time series in a vector.
2. **Propagating** obtained encodings throughout the graph with a stack of MP layers.



Full TTS model



Pros & Cons of TTS models

Pros: 😊 Easy to implement and computationally efficient.

😊 We can reuse operators we already know.


Cons: 😞 The 2-step encoding might introduce information bottlenecks.

😞 Accounting for changes in topology and dynamic edge attributes can be more problematic.

Space-then-Time


In STT approaches the two processing steps of TTS models are inverted:

1. Observations are **propagated among nodes** w.r.t. each time step using a stack of MP layers.
2. Each sequence of representations is processed by a **sequence encoder**.



☹️ They do not have the same computational advantages of TTS models.

Full STT model



Model quality assessment

Questions to answer

Consider a predictor \mathcal{F} trained to solve a time-series forecasting problem.

1. Is the predictor **optimal** for the problem at hand?
2. **Where** does the predictor appear to be sub-optimal?
3. **How** can we improve the predictor?

Remark: Multiple optimality criteria can be considered.

 Relational inductive biases can help us

Performance at task

Consider predictors $\mathcal{F}_a, \mathcal{F}_b$ from a set \mathbb{F} of models and performance metric M (e.g., MAE, MSE).

- we consider \mathcal{F}_a **better** than \mathcal{F}_b if $M(\mathcal{F}_a)$ is *statistically* better than $M(\mathcal{F}_b)$.
- we consider \mathcal{F}_a **optimal** if there is no $\mathcal{F}_b \in \mathbb{F}$ better than \mathcal{F}_a .

Can we further improve over the best model so far \mathcal{F}_a ?

- Either we **find a new model** \mathcal{F}_* better than \mathcal{F}_a
- or we need **prior knowledge** about the modeled system.

Model	M
\mathcal{F}_a	$0.145_{\pm 0.002}$
\mathcal{F}_b	$0.176_{\pm 0.005}$
\vdots	
\mathcal{F}_n	$0.158_{\pm 0.004}$
\mathcal{F}_*	$0.139_{\pm 0.001}$

Residual correlation analysis

Studying the correlation between prediction residuals $r_t^i = x_{t:t+H}^i - \hat{x}_{t:t+H}^i$ allows us for testing model optimality.

If residuals are **dependent**

⇒ there is **information** that the model **hasn't captured**

⇒ model predictions **can be improved**.

Temporal correlation

Correlation between residuals at different time steps.

Spatial correlation

Correlation between residuals at different graph nodes.

Most of the research focused on either serial correlation [15]–[17] or spatial correlation [18], [19].

Statistical tests for residual correlation

Whiteness test

H_0 : residuals are **uncorrelated** H_1 : some residuals **correlate**


Define a test statistic $C(\{\mathbf{r}_t^i\}) = C(\mathcal{F}, \{\mathbf{x}_t^i\})$ and a threshold γ such that

$$\text{If } |C(\{\mathbf{r}_t^i\})| > \gamma \implies \text{reject } H_0.$$

Remarks: Residual correlation analysis

- 😊 Is independent of specific performance measures.
- 😞 Does not quantify how much a model can improve w.r.t. a specific performance metric.
- 😊 Does not rely on comparisons with other models.

AZ-Whiteness test: a spatio-temporal test



The test is defined by statistic


$$C(\{\mathbf{r}\}) = \underbrace{\sum_t \sum_{(i,j) \in \mathcal{E}_t} w_{ijt} \operatorname{sgn}(\langle \mathbf{r}_t^i, \mathbf{r}_t^j \rangle)}_{\text{spatial edge}} + \underbrace{\sum_t \sum_i w_{it} \operatorname{sgn}(\langle \mathbf{r}_t^i, \mathbf{r}_{t+1}^i \rangle)}_{\text{temporal edge}}$$

- distribution-free and residuals can be non-identically distributed.
- computation is linear in the number of edges and time steps.


Where can we improve?

Analyzing the AZ-whiteness **test** statistic computed **on subgraphs** of the spatio-temporal graph allows for **discovering** insightful **correlation patterns**.


Spatial (or temporal) edges




Edges related to a time step



Edges related to a node



Edges related to a node



Dealing with missing data

The problem of missing data

So far, we assumed to deal with **complete sequences**, i.e., to have valid observations associated with each node (sensor) and time step.

However, time series collected by real-world sensor networks often have **missing data**, due to:


- faults, of either transient or permanent nature;
- asynchronicity among the time series;
- communication errors...

Most forecasting methods operate on complete sequences.

→ We need a way to **impute**, i.e., *reconstruct*, missing data.



Time series imputation (i)

The problem of reconstructing missing values in a sequence of data is often referred to as **time series imputation (TSI)**.



Time series imputation (ii)

We use a **mask** $m_t^i \in \{0, 1\}$ to distinguish between missing (0) and valid (1) observations.



Time series imputation


Given a window of $T \geq 1$ observations $\mathbf{X}_{<T}$ with missing values, the **time series imputation** problem consists in estimating the missing observations in the sequence

$$\mathbf{x}_t^i \sim p(\mathbf{x}_t^i | \mathcal{X}_{<T}) \quad \forall i, t \text{ such that } m_t^i = 0$$

with $\mathcal{X}_{<T} = \{\mathbf{x}_t^i | \mathbf{x}_t^i \in \mathbf{X}_{<T} \text{ and } m_t^i = 1\}$ being the **observed set**.

Deep learning for TSI

Besides standard statistical methods, deep learning approaches have become a popular alternative, in particular, **autoregressive models** (e.g., RNNs).



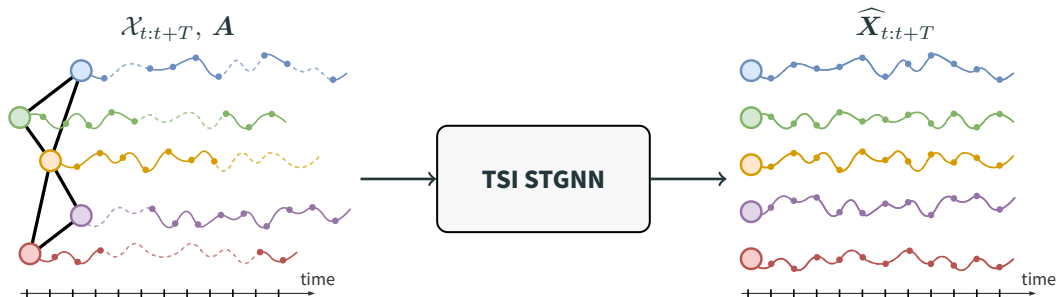
- 😊 Effective in exploiting past (and future, with bidirectional models) **node** observations...
- 😞 ...but struggle in capturing **nonlinear space-time dependencies**.

Time series imputation + relational inductive biases

Again, we can use the available relational information to condition the model, i.e.,

$$\mathbf{x}_t^i \sim p(\mathbf{x}_t^i | \mathcal{X}_{<T}, \mathbf{A})$$

As done for the forecasting problem, we can use STGNNs to address the imputation task.



Graph Recurrent Imputation Network

Cini et al. [32] propose a GCRNN that builds upon the autoregressive approach for imputation:

- A (graph-based) RNN (i.e., a GCRNN cell) is used to **encode** the sequence of **only valid observations**:

$$\mathbf{Z}_t = \text{STMP}(\mathbf{H}_{<t} \odot \mathbf{M}_{<t}, \mathcal{E}_{<t}).$$

- An additional MP layer is used as **spatial decoder**, to account for **concurrent observations at neighbors**:

$$\hat{\mathbf{x}}_t^i = \text{DEC} \left(\mathbf{z}_t^i, \text{AGGR}_{j \in \mathcal{N}(i) \setminus \{i\}} \left\{ \text{MSG}(\mathbf{z}_t^j, \mathbf{x}_t^j) \right\} \right).$$


[32] A. Cini *et al.*, “Filling the G_ap_s: Multivariate Time Series Imputation by Graph Neural Networks”, ICLR 2022.

Forecasting from Partial Observations

A different approach to the problem **avoids the reconstruction step** and considers forecasting architectures that **directly deal with irregular observations**.

The mechanisms used in imputation can be adapted to build forecasting architectures.

- 😊 Such models can be used to **jointly impute** missing observations and **forecast** future values.




Beyond imputation

Graph-based imputation methods estimate missing values at an **existing node** by using available information at **neighboring nodes**.

Question:

Can we use the same approach to **infer** observations of **virtual sensors**, i.e., fictitious nodes **not** associated with an existing sensor?




Virtual sensing

💡 Simulate the presence of a sensor by adding a node with **no data**, then let the model **infer** the corresponding time series.

Clearly, several assumptions are needed

- high degree of homogeneity of sensors,
- capability to reconstruct from observations at neighboring sensors,
- and many more...

Two virtual sensors for air quality. (from [32])



[10] I. Marisca *et al.*, “Learning to Reconstruct Missing Data from Spatiotemporal Graphs with Sparse Observations”, NeurIPS 2022.

[32] A. Cini *et al.*, “Filling the G_ap_s: Multivariate Time Series Imputation by Graph Neural Networks”, ICLR 2022.

[33] Y. Wu *et al.*, “Inductive Graph Neural Networks for Spatiotemporal Kriging”, AAAI 2021.

Coding Spatiotemporal GNNs

tsl: PyTorch Spatiotemporal Library



tsl (Torch Spatiotemporal) is a python library built upon [PyTorch](#) and [PyG](#) to accelerate research on neural spatiotemporal data processing methods, with a focus on **Graph Neural Networks**.



Notebook

Spatiotemporal Graph Neural Networks with tsl



References i

- [1] K. Benidis, S. S. Rangapuram, V. Flunkert, *et al.*, “Deep learning for time series forecasting: Tutorial and literature survey,” *ACM Comput. Surv.*, vol. 55, no. 6, 2022, issn: 0360-0300. doi: 10.1145/3533382. [Online]. Available: <https://doi.org/10.1145/3533382>.
- [2] D. Bacciu, F. Errica, A. Micheli, and M. Podda, “A gentle introduction to deep learning for graphs,” *Neural Networks*, vol. 129, pp. 203–221, 2020.
- [3] M. M. Bronstein, J. Bruna, T. Cohen, and P. Veličković, “Geometric deep learning: Grids, groups, graphs, geodesics, and gauges,” *arXiv preprint arXiv:2104.13478*, 2021.
- [4] A. Cini, I. Marisca, D. Zambon, and C. Alippi, “Graph deep learning for time series forecasting: A comprehensive methodological framework,” *arxiv*, 2023.
- [5] J. Gilmer, S. S. Schoenholz, P. F. Riley, O. Vinyals, and G. E. Dahl, “Neural message passing for quantum chemistry,” in *International conference on machine learning*, PMLR, 2017, pp. 1263–1272.
- [6] J. Chung, C. Gulcehre, K. Cho, and Y. Bengio, “Empirical evaluation of gated recurrent neural networks on sequence modeling,” *arXiv preprint arXiv:1412.3555*, 2014.

References ii

- [7] Y. Seo, M. Defferrard, P. Vandergheynst, and X. Bresson, “Structured sequence modeling with graph convolutional recurrent networks,” in *International Conference on Neural Information Processing*, Springer, 2018, pp. 362–373.
- [8] Y. Li, R. Yu, C. Shahabi, and Y. Liu, “Diffusion convolutional recurrent neural network: Data-driven traffic forecasting,” in *International Conference on Learning Representations*, 2018.
- [9] B. Yu, H. Yin, and Z. Zhu, “Spatio-temporal graph convolutional networks: A deep learning framework for traffic forecasting,” in *Proceedings of the 27th International Joint Conference on Artificial Intelligence*, 2018, pp. 3634–3640.
- [10] I. Marisca, A. Cini, and C. Alippi, “Learning to reconstruct missing data from spatiotemporal graphs with sparse observations,” in *Advances in Neural Information Processing Systems*, 2022.
- [11] Z. Wu, D. Zheng, S. Pan, Q. Gan, G. Long, and G. Karypis, “Traversenet: Unifying space and time in message passing for traffic forecasting,” *IEEE Transactions on Neural Networks and Learning Systems*, 2022.
- [12] M. Sabbaqi and E. Isufi, “Graph-time convolutional neural networks: Architecture and theoretical analysis,” *arXiv preprint arXiv:2206.15174*, 2022.

References iii

- [13] P. Montero-Manso and R. J. Hyndman, “Principles and algorithms for forecasting groups of time series: Locality and globality,” *International Journal of Forecasting*, vol. 37, no. 4, pp. 1632–1653, 2021.
- [14] A. Cini, I. Marisca, D. Zambon, and C. Alippi, “Taming local effects in graph-based spatiotemporal forecasting,” *arXiv preprint arXiv:2302.04071*, 2023.
- [15] JRM Hosking, “Equivalent Forms of the Multivariate Portmanteau Statistic,” *Journal of the Royal Statistical Society: Series B (Methodological)*, vol. 43, no. 2, pp. 261–262, 1981.
- [16] Z. Li, C. Lam, J. Yao, and Q. Yao, “On Testing for High-Dimensional White Noise,” *The Annals of Statistics*, vol. 47, no. 6, pp. 3382–3412, 2019.
- [17] A. Bose and W. Hachem, “A Whiteness Test Based on the Spectral Measure of Large Non-Hermitian Random Matrices,” in *ICASSP 2020-2020 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, IEEE, 2020, pp. 8768–8771.
- [18] P. A. P. Moran, “Notes on Continuous Stochastic Phenomena,” *Biometrika*, vol. 37, no. 1/2, pp. 17–23, 1950, ISSN: 0006-3444. DOI: 10.2307/2332142.

References iv

- [19] A. D. Cliff and K. Ord, “Spatial Autocorrelation: A Review of Existing and New Measures with Applications,” *Economic Geography*, vol. 46, pp. 269–292, 1970, ISSN: 0013-0095. DOI: 10.2307/143144.
- [20] D. Zambon and C. Alippi, “AZ-whiteness test: A test for signal uncorrelation on spatio-temporal graphs,” in *Advances in Neural Information Processing Systems*, A. H. Oh, A. Agarwal, D. Belgrave, and K. Cho, Eds., 2022.
- [21] D. Zambon and C. Alippi, “Where and how to improve graph-based spatio-temporal predictors,” *arXiv preprint arXiv:2302.01701*, 2023.
- [22] A. Cini, D. Zambon, and C. Alippi, “Sparse graph learning from spatiotemporal time series,” *Journal of Machine Learning Research*, vol. 24, no. 242, pp. 1–36, 2023.
- [23] Z. Wu, S. Pan, G. Long, J. Jiang, and C. Zhang, “Graph wavenet for deep spatial-temporal graph modeling,” in *Proceedings of the 28th International Joint Conference on Artificial Intelligence*, 2019, pp. 1907–1913.
- [24] T. Kipf, E. Fetaya, K.-C. Wang, M. Welling, and R. Zemel, “Neural relational inference for interacting systems,” in *International conference on machine learning*, PMLR, 2018, pp. 2688–2697.

References v

- [25] G. Ditzler, M. Roveri, C. Alippi, and R. Polikar, “Learning in nonstationary environments: A survey,” *IEEE Computational Intelligence Magazine*, vol. 10, no. 4, pp. 12–25, 2015.
- [26] G. Panagopoulos, G. Nikolentzos, and M. Vazirgiannis, “Transfer graph neural networks for pandemic forecasting,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 35, 2021, pp. 4838–4845.
- [27] T. Mallick, P. Balaprakash, E. Rask, and J. Macfarlane, “Transfer learning with graph neural networks for short-term highway traffic forecasting,” in *2020 25th International Conference on Pattern Recognition (ICPR)*, IEEE, 2021, pp. 10 367–10 374.
- [28] W. Hamilton, Z. Ying, and J. Leskovec, “Inductive representation learning on large graphs,” *Advances in neural information processing systems*, vol. 30, 2017.
- [29] Y. Rong, W. Huang, T. Xu, and J. Huang, “Dropedge: Towards deep graph convolutional networks on node classification,” in *International Conference on Learning Representations*, 2020.
- [30] F. Frasca, E. Rossi, D. Eynard, B. Chamberlain, M. Bronstein, and F. Monti, “SIGN: Scalable inception graph neural networks,” *arXiv preprint arXiv:2004.11198*, 2020.

References vi

- [31] A. Cini, I. Marisca, F. M. Bianchi, and C. Alippi, “Scalable spatiotemporal graph neural networks,” *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 37, no. 6, pp. 7218–7226, Jun. 2023. doi: 10.1609/aaai.v37i6.25880.
- [32] A. Cini, I. Marisca, and C. Alippi, “Filling the g_ap_s: Multivariate time series imputation by graph neural networks,” in *International Conference on Learning Representations*, 2022. [Online]. Available: <https://openreview.net/forum?id=k0u3-S3wJ7>.
- [33] Y. Wu, D. Zhuang, A. Labbe, and L. Sun, “Inductive Graph Neural Networks for Spatiotemporal Kriging,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 35, 2021, pp. 4478–4485.